

```
# Import required libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Optional: To display charts inline
%matplotlib inline

#Load the datasets
purchase_df = pd.read_csv('QVI_purchase_behaviour.csv')
transaction_df = pd.read_excel('QVI_transaction_data.xlsx')

#check missing values
print(purchase_df.info())
print(transaction_df.info())
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 72637 entries, 0 to 72636
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   LYLTY_CARD_NBR        72637 non-null  int64
1   LIFESTAGE             72637 non-null  object
2   PREMIUM_CUSTOMER     72637 non-null  object
dtypes: int64(1), object(2)
memory usage: 1.7+ MB
None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 264839 entries, 0 to 264838
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   DATE                  264836 non-null  datetime64[ns]
1   STORE_NBR             264836 non-null  float64
2   LYLTY_CARD_NBR        264836 non-null  float64
3   TXN_ID                264836 non-null  float64
4   PROD_NBR              264836 non-null  float64
5   Update_Product_Name   264836 non-null  object
6   PROD_QTY              264838 non-null  float64
7   TOT_SALES             264838 non-null  float64
dtypes: datetime64[ns](1), float64(6), object(1)
memory usage: 16.2+ MB
None
```

```
#Description Statistics
print(transaction_df.describe())
transaction_df.isnull().sum()
```



	DATE	STORE_NBR	LYLTY_CARD_NBR	\
count	264836	264836.00000	2.648360e+05	
mean	2018-12-30 00:52:12.879215360	135.08011	1.355495e+05	
min	2018-07-01 00:00:00	1.00000	1.000000e+03	
25%	2018-09-30 00:00:00	70.00000	7.002100e+04	
50%	2018-12-30 00:00:00	130.00000	1.303575e+05	
75%	2019-03-31 00:00:00	203.00000	2.030942e+05	
max	2019-06-30 00:00:00	272.00000	2.373711e+06	
std	NaN	76.78418	8.057998e+04	

	TXN_ID	PROD_NBR	PROD_QTY	TOT_SALES
count	2.648360e+05	264836.000000	264838.000000	264838.000000
mean	1.351583e+05	56.583157	1.907324	7.304259
min	1.000000e+00	1.000000	1.000000	1.500000
25%	6.760150e+04	28.000000	2.000000	5.400000
50%	1.351375e+05	56.000000	2.000000	7.400000
75%	2.027012e+05	85.000000	2.000000	9.200000
max	2.415841e+06	114.000000	200.000000	650.000000
std	7.813303e+04	32.826638	0.643683	3.083303

0

DATE	3
STORE_NBR	3
LYLTY_CARD_NBR	3
TXN_ID	3
PROD_NBR	3
Update_Product_Name	3
PROD_QTY	1
TOT_SALES	1

dtype: int64

#merge datasets

```
merged_df = pd.merge(transaction_df, purchase_df, on='LYLTY_CARD_NBR', how='left')
print(merged_df)
```



	DATE	STORE_NBR	LYLTY_CARD_NBR	TXN_ID	PROD_NBR	\
0	2018-07-01	4.0	4074.0	2982.0	57.0	
1	2018-07-01	4.0	4149.0	3333.0	16.0	
2	2018-07-01	4.0	4196.0	3539.0	24.0	
3	2018-07-01	5.0	5026.0	4525.0	42.0	
4	2018-07-01	7.0	7215.0	7176.0	16.0	
...	
264834	2019-06-22	208.0	208205.0	207318.0	37.0	
264835	2019-06-22	264.0	264149.0	262909.0	25.0	
264836	NaT	NaN	NaN	NaN	NaN	
264837	NaT	NaN	NaN	NaN	NaN	

264838

NaT

NaN

NaN

NaN


NaN

	Update_Product_Name	PROD_QTY	TOT_SALES	\
0	Old El Paso Salsa Dip Tomato Mild 300g	1.0	5.1	
1	Smiths Crinkle Chips Salt & Vinegar 330g	1.0	5.7	
2	Grain Waves Sweet Chilli 210g	1.0	3.6	
3	Doritos Corn Chip Mexican Jalapeno 150g	1.0	3.9	
4	Smiths Crinkle Chips Salt & Vinegar 330g	1.0	5.7	
...	
264834	Smiths Thinly Swt Chli&S/Cream175G	5.0	15.0	
264835	Pringles SourCream Onion 134g	5.0	18.5	
264836	NaN	NaN	NaN	
264837	NaN	5.0	18.5	
264838	NaN	3.0	11.8	

	LIFESTAGE	PREMIUM_CUSTOMER
0	MIDAGE SINGLES/COUPLES	Budget
1	MIDAGE SINGLES/COUPLES	Budget
2	MIDAGE SINGLES/COUPLES	Budget
3	MIDAGE SINGLES/COUPLES	Budget
4	MIDAGE SINGLES/COUPLES	Budget
...
264834	YOUNG SINGLES/COUPLES	Premium
264835	YOUNG SINGLES/COUPLES	Premium
264836	NaN	NaN
264837	NaN	NaN
264838	NaN	NaN

[264839 rows x 10 columns]

```
transaction_df.dropna(inplace=True)
transaction_df.isnull().sum()
```



	0
DATE	0
STORE_NBR	0
LYLTY_CARD_NBR	0
TXN_ID	0
PROD_NBR	0
Update_Product_Name	0
PROD_QTY	0
TOT_SALES	0

dtype: int64

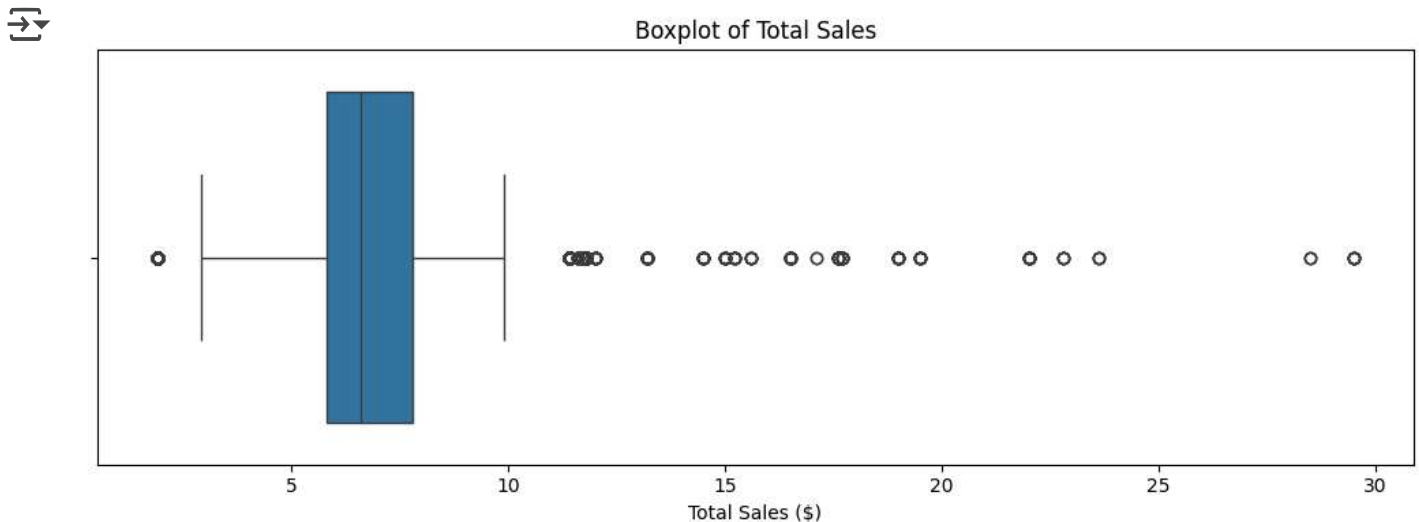
```
#Keep only chip-related products
merged_df = merged_df[merged_df['Update_Product_Name'].str.contains("chip|chips", case=False)]

# Remove outlier customer who bought 200 units
merged_df = merged_df[merged_df['LYLTY_CARD_NBR'] != 226000]

# Check for missing dates
merged_df['DATE'] = pd.to_datetime(merged_df['DATE'])
full_dates = pd.date_range(start="2018-07-01", end="2019-06-30")
actual_dates = merged_df['DATE'].dt.normalize().unique()
missing_dates = [date for date in full_dates if date not in actual_dates]
print("Missing dates:", missing_dates)
```

```
➞ Missing dates: [Timestamp('2018-12-25 00:00:00')]
```

```
# Boxplot of sales to check outliers
plt.figure(figsize=(10, 4))
sns.boxplot(x=merged_df['TOT_SALES'])
plt.title("Boxplot of Total Sales")
plt.xlabel("Total Sales ($)")
plt.tight_layout()
plt.show()
```



```
# Remove outliers using IQR
Q1 = merged_df['TOT_SALES'].quantile(0.25)
Q3 = merged_df['TOT_SALES'].quantile(0.75)
IQR = Q3 - Q1
merged_df = merged_df[(merged_df['TOT_SALES'] >= Q1 - 1.5 * IQR) &
                      (merged_df['TOT_SALES'] <= Q3 + 1.5 * IQR)]
```

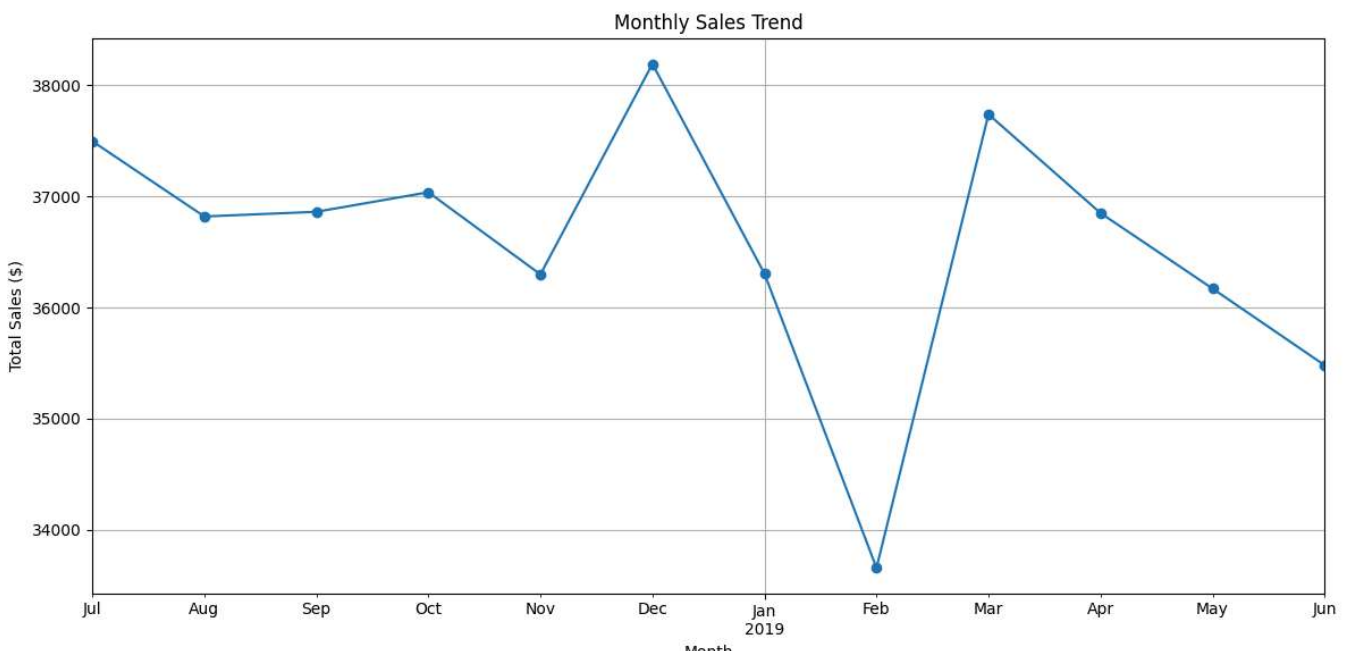
```
# Extract PACK_SIZE from PROD_NAME
merged_df['PACK_SIZE'] = merged_df['Update_Product_Name'].str.extract(r'(\d+)').astype(float)

# Extract and clean BRAND
merged_df['BRAND'] = merged_df['Update_Product_Name'].str.extract(r'^([A-Za-z]+)')
brand_fix = {'RED': 'RRD', 'SNBTS': 'SUNBITES', 'WW': 'WOOLWORTHS', 'SMITH': 'SMITHS',
             'NCC': 'NATURAL', 'DORITO': 'DORITOS', 'INFZNS': 'INFUZIONI', 'GRAIN': 'GRNWVES'}
merged_df['BRAND'].replace(brand_fix, inplace=True)
```


ed1ba893524e>:5: FutureWarning: A value is trying to be set on a copy of a DataFrame or change in pandas 3.0. This inplace method will never work because the intermediate object is a view. To avoid this, use `df[col].method(value, inplace=True)`, try using `df.method({col: value}, inplace=True)`.

```
# Monthly sales trend
monthly_sales = merged_df.resample('M', on='DATE')['TOT_SALES'].sum()
monthly_sales.plot(kind='line', marker='o', figsize=(12, 6), title='Monthly Sales Trend')
plt.ylabel('Total Sales ($)')
plt.xlabel('Month')
plt.grid(True)
plt.tight_layout()
plt.show()
```

<ipython-input-15-5a280820b557>:2: FutureWarning: 'M' is deprecated and will be removed in a future version. Use 'month' instead.




```
# Total sales by customer segment
sales_by_type = (merged_df.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['TOT_SALES']
                  .sum()
                  .reset_index()
                  .sort_values(by='TOT_SALES', ascending=False))
print(sales_by_type)
```



	LIFESTAGE	PREMIUM_CUSTOMER	TOT_SALES
6	OLDER FAMILIES	Budget	38481.0
13	RETIREEES	Mainstream	35074.7
19	YOUNG SINGLES/COUPLES	Mainstream	33518.9
15	YOUNG FAMILIES	Budget	31905.7
10	OLDER SINGLES/COUPLES	Mainstream	30728.1
9	OLDER SINGLES/COUPLES	Budget	30615.1
11	OLDER SINGLES/COUPLES	Premium	29595.4
12	RETIREEES	Budget	25493.5
7	OLDER FAMILIES	Mainstream	24268.3
16	YOUNG FAMILIES	Mainstream	21895.2
14	RETIREEES	Premium	21268.2
1	MIDAGE SINGLES/COUPLES	Mainstream	19857.4
17	YOUNG FAMILIES	Premium	19736.2
8	OLDER FAMILIES	Premium	18352.9
18	YOUNG SINGLES/COUPLES	Budget	14735.6
2	MIDAGE SINGLES/COUPLES	Premium	13317.3
20	YOUNG SINGLES/COUPLES	Premium	10088.2
0	MIDAGE SINGLES/COUPLES	Budget	8613.0
3	NEW FAMILIES	Budget	4999.5
4	NEW FAMILIES	Mainstream	3757.1
5	NEW FAMILIES	Premium	2595.9

```
# 📊 Bar chart of customer segment sales
plt.figure(figsize=(10, 6))
sns.barplot(data=sales_by_type, x='LIFESTAGE', y='TOT_SALES', hue='PREMIUM_CUSTOMER')
plt.title('Sales by Customer Segment')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
# Average units per customer
avg_units = merged_df.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['PROD_QTY'].sum() / \
    merged_df.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['LYLTY_CARD_NBR'].nunique()
print("Average units per customer:\n", avg_units)
```



Average units per customer:

LIFESTAGE	PREMIUM_CUSTOMER	
MIDAGE SINGLES/COUPLES	Budget	3.020095
	Mainstream	3.031686
	Premium	3.026134
NEW FAMILIES	Budget	2.636531
	Mainstream	2.527316
	Premium	2.537415
OLDER FAMILIES	Budget	3.810909
	Mainstream	3.869010
	Premium	3.843923
OLDER SINGLES/COUPLES	Budget	3.041424
	Mainstream	3.126477
	Premium	3.077309
RETIRES	Budget	2.879181
	Mainstream	2.856139
	Premium	2.848951
YOUNG FAMILIES	Budget	3.718640
	Mainstream	3.690355
	Premium	3.722716
YOUNG SINGLES/COUPLES	Budget	2.507090
	Mainstream	2.542340
	Premium	2.543024

dtype: float64

```
# 💰 Average price per unit
avg_price = merged_df.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['TOT_SALES'].sum() / \
    merged_df.groupby(['LIFESTAGE', 'PREMIUM_CUSTOMER'])['PROD_QTY'].sum()
print("Average price per unit:\n", avg_price)
```

```
➡ Average price per unit:
  LIFESTAGE      PREMIUM_CUSTOMER
MIDAGE SINGLES/COUPLES Budget      3.371037
                        Mainstream  3.517697
                        Premium     3.382601
NEW FAMILIES      Budget      3.498600
                        Mainstream  3.531109
                        Premium     3.479759
OLDER FAMILIES    Budget      3.338046
                        Mainstream  3.339981
                        Premium     3.297323
OLDER SINGLES/COUPLES Budget      3.446094
                        Mainstream  3.414992
                        Premium     3.442126
RETIREEES         Budget      3.484623
                        Mainstream  3.450536
                        Premium     3.480314
YOUNG FAMILIES    Budget      3.352848
                        Mainstream  3.346355
                        Premium     3.363934
YOUNG SINGLES/COUPLES Budget      3.333846
                        Mainstream  3.613118
                        Premium     3.314126

dtype: float64
```

```
# 🎯 Brand affinity: Mainstream + Young Singles/Couples
segment = merged_df[(merged_df['LIFESTAGE'] == "YOUNG SINGLES/COUPLES") &
                    (merged_df['PREMIUM_CUSTOMER'] == "Mainstream")]
segment_brand_share = segment['BRAND'].value_counts(normalize=True)
other_brand_share = merged_df[~merged_df.index.isin(segment.index)]['BRAND'].value_counts(normalize=True)
brand_affinity = (segment_brand_share / other_brand_share).dropna().sort_values(ascending=False)
print("Brand Affinity:\n", brand_affinity)
```

```
➡ Brand Affinity:
  BRAND
Doritos      1.286385
Cobs          1.217147
Tostitos     1.213476
Thins        1.125839
French       0.723629
Smiths       0.713365
Natural      0.699358
WOOLWORTHS   0.496231
Name: proportion, dtype: float64
```



```
# Pack size affinity
segment_qty = segment['PROD_QTY'].sum()
other = merged_df[~merged_df.index.isin(segment.index)]
other_qty = other['PROD_QTY'].sum()
segment_pack_share = segment.groupby('PACK_SIZE')['PROD_QTY'].sum() / segment_qty
other_pack_share = other.groupby('PACK_SIZE')['PROD_QTY'].sum() / other_qty
pack_affinity = (segment_pack_share / other_pack_share).dropna().sort_values(ascending=False)
print("Pack Size Affinity:\n", pack_affinity)
```

➞ Pack Size Affinity:

PACK_SIZE	
380.0	2.319916
330.0	2.247313
110.0	1.227404
150.0	1.213443
170.0	1.071249
175.0	0.948822
160.0	0.532697
200.0	0.494835

Name: PROD_QTY, dtype: float64

```
merged_df.to_csv("QVI_cleaned_data.csv", index=False)
```

Start coding or [generate](#) with AI.